

Data acquisition

In [1]:

```
# library for importing datasets

import pandas as pd
```

In [2]:

```
# movies dataset

movies_data = pd.read_table(r"C:\Users\PKN\Data_Science\DS_w_Python\Final Projects\Proj
```

C:\Users\PKN\anaconda3\lib\site-packages\ipykernel_launcher.py:3: ParserWarning: Falling back to the 'python' engine because the 'c' engine does not support regex separators (separators > 1 char and different from '\s+' are interpreted as regex); you can avoid this warning by specifying engine='python'.

This is separate from the ipykernel package so we can avoid doing imports until

In [3]:

```
movies_data.shape
```

Out[3]: (3883, 3)

In [4]:

```
movies_data.head()
```

Out[4]:

	MovieID	Title	Genres
0	1	Toy Story (1995)	Animation Children's Comedy
1	2	Jumanji (1995)	Adventure Children's Fantasy
2	3	Grumpier Old Men (1995)	Comedy Romance
3	4	Waiting to Exhale (1995)	Comedy Drama
4	5	Father of the Bride Part II (1995)	Comedy

In [5]:

```
# ratings dataset

ratings_data = pd.read_table(r"C:\Users\PKN\Data_Science\DS_w_Python\Final Projects\Pro
```

C:\Users\PKN\anaconda3\lib\site-packages\ipykernel_launcher.py:3: ParserWarning: Falling back to the 'python' engine because the 'c' engine does not support regex separators (separators > 1 char and different from '\s+' are interpreted as regex); you can avoid this warning by specifying engine='python'.

This is separate from the ipykernel package so we can avoid doing imports until

In [6]:

```
ratings_data.shape
```

Out[6]: (1000209, 4)

In [7]:

```
ratings_data.head()
```

```
Out[7]:
```

	UserID	MovieID	Ratings	TimeStamp
0	1	1193	5	978300760
1	1	661	3	978302109
2	1	914	3	978301968
3	1	3408	4	978300275
4	1	2355	5	978824291

```
In [8]: # ratings dataset

users_data = pd.read_table(r"C:\Users\PKN\Data_Science\DS_w_Python\Final Projects\Proje
```

C:\Users\PKN\anaconda3\lib\site-packages\ipykernel_launcher.py:3: ParserWarning: Falling back to the 'python' engine because the 'c' engine does not support regex separators (separators > 1 char and different from '\s+' are interpreted as regex); you can avoid this warning by specifying engine='python'.

This is separate from the ipykernel package so we can avoid doing imports until

```
In [9]: users_data.shape
```

```
Out[9]: (6040, 5)
```

```
In [10]: users_data.head()
```

```
Out[10]:
```

	UserID	Gender	Age	Occupation	Zip-Code
0	1	F	1	10	48067
1	2	M	56	16	70072
2	3	M	25	15	55117
3	4	M	45	7	02460
4	5	M	25	20	55455

Creating a new Master Dataset by merging the previous 3 datasets

```
In [11]: movielens = movies_data.merge(ratings_data, on = 'MovieID', how = 'inner').merge(users_
```

```
In [12]: movielens.shape
```

```
Out[12]: (1000209, 10)
```

```
In [13]: movielens.head()
```

```
Out[13]:
```

	MovieID	Title	Genres	UserID	Ratings	TimeStamp	Gender	Age
--	---------	-------	--------	--------	---------	-----------	--------	-----

	MovieID	Title	Genres	UserID	Ratings	TimeStamp	Gender	Age
0	1	Toy Story (1995)	Animation Children's Comedy	1	5	978824268	F	1
1	48	Pocahontas (1995)	Animation Children's Musical Romance	1	5	978824351	F	1
2	150	Apollo 13 (1995)	Drama	1	5	978301777	F	1
3	260	Star Wars: Episode IV - A New Hope (1977)	Action Adventure Fantasy Sci-Fi	1	4	978300760	F	1
4	527	Schindler's List (1993)	Drama War	1	5	978824195	F	1

Exploratory Data Analysis (EDA)

Visualize user age distribution

```
In [14]: # Users with Different Age Groups

movielens['Age'].value_counts()
```


```
Out[14]: 25    395556
        35    199003
        18    183536
        45     83633
        50     72490
        56     38780
         1     27211
        Name: Age, dtype: int64
```

```
In [15]: # Libraries for plotting

import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [16]: movielens['Age'].value_counts().plot(kind='bar')
        plt.title('Age distribution chart')
        plt.xlabel('Age')
        plt.ylabel('User Counts')
```

```
Out[16]: Text(0, 0.5, 'User Counts')
```



```
In [17]: # -----EXTRA-----
# -----Another style of plotting-----
# Libraries for plotting

import seaborn as sns
```


```
In [18]: # -----EXTRA-----

sns.countplot(movielens['Age'])
```

C:\Users\PKN\anaconda3\lib\site-packages\seaborn_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning

```
Out[18]: <AxesSubplot:xlabel='Age', ylabel='count'>
```



```
In [19]: # -----EXTRA-----


sns.countplot(movielens['Occupation'])
```

C:\Users\PKN\anaconda3\lib\site-packages\seaborn_decorators.py:43: FutureWarning: Pass

the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning


Out[19]: <AxesSubplot:xlabel='Occupation', ylabel='count'>



In [20]:

```
# -----EXTRA-----  
# checking correletion between 'Age' and 'Occupation'  
  
slice_moviels = pd.DataFrame(moviels[['Age', 'Occupation']])  
sns.boxplot(data = slice_moviels)
```

Out[20]: <AxesSubplot:>




In [21]:

```
# -----EXTRA-----  
  
sns.distplot(moviels['Age'], bins=8)
```

C:\Users\PKN\anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)

Out[21]: <AxesSubplot:xlabel='Age', ylabel='Density'>



Visualize overall rating by users

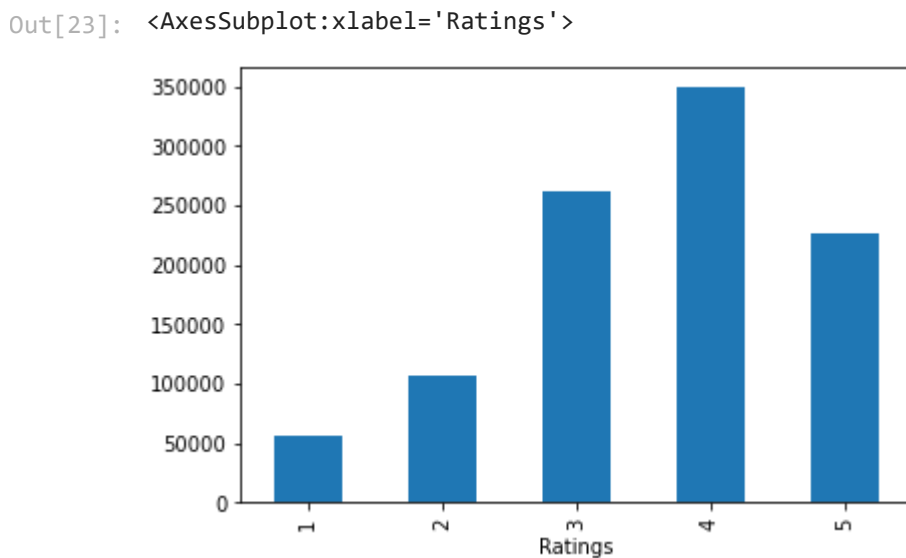
```
In [22]: # Check 'Ratings' Distribution

movielens.groupby('Ratings', axis=0).UserID.count()
```

```
Out[22]: Ratings
1      56174
2     107557
3     261197
4     348971
5     226310
Name: UserID, dtype: int64
```

```
In [23]: # Visualize

movielens.groupby('Ratings', axis=0).UserID.count().plot(kind='bar')
```




```
In [24]: # -----EXTRA-----

sns.countplot(movielens['Ratings'])
```

```
C:\Users\PKN\anaconda3\lib\site-packages\seaborn\_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
FutureWarning
```

Out[24]: <AxesSubplot:xlabel='Ratings', ylabel='count'>



visualize the user rating of the movie 'Toy Story'

```
In [25]: toystoryRating = movielens[movielens['Title'].str.contains('Toy Story') == True]
toystoryRating
```

Out[25]:

	MovieID	Title	Genres	UserID	Ratings	TimeStamp	Gender	Age	Occ
0	1	Toy Story (1995)	Animation Children's Comedy	1	5	978824268	F	1	
50	3114	Toy Story 2 (1999)	Animation Children's Comedy	1	4	978302174	F	1	
53	1	Toy Story (1995)	Animation Children's Comedy	6	4	978237008	F	50	
124	1	Toy Story (1995)	Animation Children's Comedy	8	4	978233496	M	25	
263	1	Toy Story (1995)	Animation Children's Comedy	9	5	978225952	M	25	
...
998988	3114	Toy Story 2 (1999)	Animation Children's Comedy	3023	4	970471948	F	25	

	MovieID	Title	Genres	UserID	Ratings	TimeStamp	Gender	Age	Occ
999027	3114	Toy Story 2 (1999)	Animation Children's Comedy	5800	5	958015250	M	35	
999486	3114	Toy Story 2 (1999)	Animation Children's Comedy	2189	4	974607816	M	1	
999869	3114	Toy Story 2 (1999)	Animation Children's Comedy	159	4	989966944	F	45	
1000192	3114	Toy Story 2 (1999)	Animation Children's Comedy	5727	5	958492554	M	25	

3662 rows × 10 columns




```
In [26]: toystoryRating.groupby(["Title", "Ratings"]).size()
```

```
Out[26]: Title
Toy Story (1995)    1      16
                  2      61
                  3     345
                  4     835
                  5     820
Toy Story 2 (1999) 1      25
                  2      44
                  3     214
                  4     578
                  5     724
dtype: int64
```

```
In [27]: # Visualize
sns.countplot(toystoryRating['Ratings'])
```

C:\Users\PKN\anaconda3\lib\site-packages\seaborn_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
FutureWarning


```
Out[27]: <AxesSubplot:xlabel='Ratings', ylabel='count'>
```


```
In [28]: sns.swarmplot(x = 'Gender', y = 'Ratings', data = toystoryRating)
```

C:\Users\PKN\anaconda3\lib\site-packages\seaborn\categorical.py:1296: UserWarning: 89.1% of the points cannot be placed; you may want to decrease the size of the markers or use stripplot.
 warnings.warn(msg, UserWarning)
 C:\Users\PKN\anaconda3\lib\site-packages\seaborn\categorical.py:1296: UserWarning: 95.3% of the points cannot be placed; you may want to decrease the size of the markers or use stripplot.
 warnings.warn(msg, UserWarning)

```
Out[28]: <AxesSubplot:xlabel='Gender', ylabel='Ratings'>
```




```
In [29]: toystoryRating.groupby(["Title", "Ratings"]).size().unstack().plot(kind='barh', stacked=F, plt.show())
```



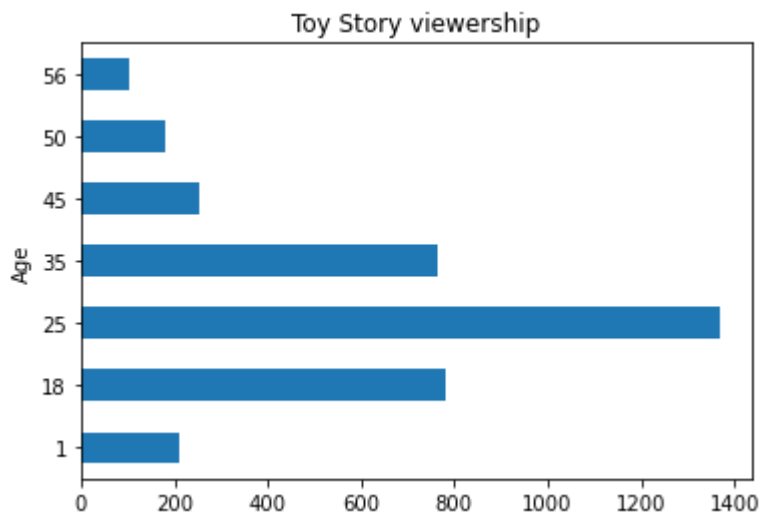
visualize the viewership of the movie "Toy Story" by age group

```
In [30]: toystoryRating.groupby('Age', axis = 0).Ratings.count().plot(kind = 'bar')
```

```
Out[30]: <AxesSubplot:xlabel='Age'>
```



```
In [31]: toystoryRating.groupby(["Age"]).size().plot(kind='barh')
plt.title('Toy Story viewership')
plt.show()
```




visualize the top 25 movies by viewership

```
In [32]: top25Movies = movielens.groupby('Title').size().sort_values(ascending=False)[:25]
         top25Movies
```

```
Out[32]: Title
American Beauty (1999)          3428
Star Wars: Episode IV - A New Hope (1977)    2991
Star Wars: Episode V - The Empire Strikes Back (1980)  2990
Star Wars: Episode VI - Return of the Jedi (1983)    2883
Jurassic Park (1993)             2672
Saving Private Ryan (1998)        2653
Terminator 2: Judgment Day (1991)  2649
Matrix, The (1999)               2590
Back to the Future (1985)        2583
Silence of the Lambs, The (1991)  2578
Men in Black (1997)              2538
Raiders of the Lost Ark (1981)    2514
 Fargo (1996)                    2513
Sixth Sense, The (1999)          2459
Braveheart (1995)                2443
Shakespeare in Love (1998)       2369
Princess Bride, The (1987)       2318
Schindler's List (1993)          2304
L.A. Confidential (1997)         2288
Groundhog Day (1993)             2278
E.T. the Extra-Terrestrial (1982)  2269
Star Wars: Episode I - The Phantom Menace (1999)    2250
Being John Malkovich (1999)      2241
Shawshank Redemption, The (1994)  2227
Godfather, The (1972)            2223
dtype: int64
```

```
In [33]: top25Movies.plot(kind='barh',alpha=0.6,figsize=(20,10))
         plt.xlabel("Viewership Ratings Count")
         plt.ylabel("Movies (Top 25)")
         plt.title("Top 25 movies by viewership rating")
         plt.show()
```



rating for a particular user of user id = 2696

In [34]:

```
id2696Rating = movielens[movielens.UserID == 2696]
id2696Rating
```

Out[34]:

	MovieID	Title	Genres	UserID	Ratings	TimeStamp	Gender	Age
991035	350	Client, The (1994)	Drama Mystery Thriller	2696	3	973308886	M	25
991036	800	Lone Star (1996)	Drama Mystery	2696	5	973308842	M	25
991037	1092	Basic Instinct (1992)	Mystery Thriller	2696	4	973308886	M	25
991038	1097	E.T. the Extra-Terrestrial (1982)	Children's Drama Fantasy Sci-Fi	2696	3	973308690	M	25
991039	1258	Shining, The (1980)	Horror	2696	4	973308710	M	25
991040	1270	Back to the Future (1985)	Comedy Sci-Fi	2696	2	973308676	M	25
991041	1589	Cop Land (1997)	Crime Drama Mystery	2696	3	973308865	M	25
991042	1617	L.A. Confidential (1997)	Crime Film-Noir Mystery Thriller	2696	4	973308842	M	25
991043	1625	Game, The (1997)	Mystery Thriller	2696	4	973308842	M	25

	MovieID	Title	Genres	UserID	Ratings	TimeStamp	Gender	Age
991044	1644	I Know What You Did Last Summer (1997)	Horror Mystery Thriller	2696	2	973308920	M	25
991045	1645	Devil's Advocate, The (1997)	Crime Horror Mystery Thriller	2696	4	973308904	M	25
991046	1711	Midnight in the Garden of Good and Evil (1997)	Comedy Crime Drama Mystery	2696	4	973308904	M	25
991047	1783	Palmetto (1998)	Film-Noir Mystery Thriller	2696	4	973308865	M	25
991048	1805	Wild Things (1998)	Crime Drama Mystery Thriller	2696	4	973308886	M	25
991049	1892	Perfect Murder, A (1998)	Mystery Thriller	2696	4	973308904	M	25
991050	2338	I Still Know What You Did Last Summer (1998)	Horror Mystery Thriller	2696	2	973308920	M	25
991051	2389	Psycho (1998)	Crime Horror Thriller	2696	4	973308710	M	25
991052	2713	Lake Placid (1999)	Horror Thriller	2696	1	973308710	M	25
991053	3176	Talented Mr. Ripley, The (1999)	Drama Mystery Thriller	2696	4	973308865	M	25
991054	3386	JFK (1991)	Drama Mystery	2696	1	973308842	M	25


Visualize the rating data by user of user if = 2696

In [35]:

```
sns.countplot(id2696Rating['Ratings'])
```


C:\Users\PKN\anaconda3\lib\site-packages\seaborn_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
FutureWarning

Out[35]: <AxesSubplot:xlabel='Ratings', ylabel='count'>



```
In [36]: movielens[movielens.UserID == 2696].groupby('Ratings').Ratings.count().plot(kind='bar')
```

```
Out[36]: <AxesSubplot:xlabel='Ratings'>
```




```
In [37]: sliceData = pd.DataFrame(movielens.iloc[0:10000, :])
```

```
In [38]: sliceData.columns
```

```
Out[38]: Index(['MovieID', 'Title', 'Genres', 'UserID', 'Ratings', 'TimeStamp',
               'Gender', 'Age', 'Occupation', 'Zip-Code'],
              dtype='object')
```

```
In [39]: sns.stripplot(y = 'Genres', x = 'Occupation', data = sliceData)
```

```
Out[39]: <AxesSubplot:xlabel='Occupation', ylabel='Genres'>
```



Perform Machine Learning

first 500 extracted records on MovieID, Age, Occupation with Rating as label

```
In [40]: movielens.columns
```

```
Out[40]: Index(['MovieID', 'Title', 'Genres', 'UserID', 'Ratings', 'TimeStamp',  
              'Gender', 'Age', 'Occupation', 'Zip-Code'],  
              dtype='object')
```

```
In [41]: first500 = movielens.iloc[:500, [0,7,8,4]]
```

```
In [42]: first500.shape
```

```
Out[42]: (500, 4)
```

```
In [43]: first500.head()
```

```
Out[43]:
```

	MovieID	Age	Occupation	Ratings
0	1	1	10	5
1	48	1	10	5
2	150	1	10	5
3	260	1	10	4
4	527	1	10	5

```
In [44]: fetaures = first500.iloc[:,[0,1,2]].values  
label = first500.iloc[:, -1].values
```

Create train and test data

```
In [45]:
```

```
#import Library for ML model

from sklearn.model_selection import train_test_split
```

```
In [46]: x_train, x_test, y_train, y_test = train_test_split(fetaures, label, test_size=0.2, ran
```

Let's create the ML Model for prediction

```
In [47]: # import KNN model Library

from sklearn.neighbors import KNeighborsClassifier
```

```
In [48]: knnmodel = KNeighborsClassifier(n_neighbors = 7)
knnmodel.fit(x_train, y_train)
```

```
Out[48]: KNeighborsClassifier(n_neighbors=7)
```

```
In [49]: print('Accuracy on training data:', knnmodel.score(x_train, y_train))
```

Accuracy on training data: 0.5125

```
In [50]: print('Accuracy on test data:', knnmodel.score(x_test, y_test))
```

Accuracy on test data: 0.35

*The accuracy score is **very bad** because we are feeding only 500 data points*


```
In [51]: y_pred = knnmodel.predict(x_test)
y_pred
```

```
Out[51]: array([4, 4, 4, 5, 4, 5, 4, 4, 3, 4, 4, 4, 4, 4, 5, 5, 3, 4, 3, 3, 4, 4,
               4, 3, 4, 4, 4, 4, 3, 3, 4, 3, 4, 4, 5, 4, 5, 5, 4, 3, 4, 4, 4, 3,
               4, 4, 4, 4, 5, 3, 4, 4, 5, 4, 4, 4, 4, 5, 4, 4, 4, 3, 4, 5, 4, 3,
               4, 3, 5, 4, 3, 5, 3, 4, 5, 4, 4, 3, 4, 5, 4, 3, 4, 3, 3, 4, 5, 3,
               3, 4, 3, 4, 3, 4, 5, 4, 5, 5, 4, 4], dtype=int64)
```

MovieID Plot

```
In [52]: plt.hist(first500.MovieID)
```


```
Out[52]: (array([98., 76., 90., 49., 45., 32., 26., 27., 38., 19.]),
          array([1.0000e+00, 3.9570e+02, 7.9040e+02, 1.1851e+03, 1.5798e+03,
                1.9745e+03, 2.3692e+03, 2.7639e+03, 3.1586e+03, 3.5533e+03,
                3.9480e+03]),
          <BarContainer object of 10 artists>)
```

Age Plot

```
In [53]: plt.hist(first500.Age)
```


```
Out[53]: (array([ 53.,   0.,   0.,   0., 245.,   0., 131.,   0.,   0.,  71.]),
          array([ 1. ,  5.9, 10.8, 15.7, 20.6, 25.5, 30.4, 35.3, 40.2, 45.1, 50. ]),
          <BarContainer object of 10 artists>)
```



Occupation Plot

```
In [54]: plt.hist(first500.Occupation)
```

```
Out[54]: (array([131.,   0.,   0.,   0.,   0., 124., 139.,   0.,   0., 106.]),
          array([ 1. ,  2.6,  4.2,  5.8,  7.4,  9. , 10.6, 12.2, 13.8, 15.4, 17. ]),
          <BarContainer object of 10 artists>)
```



```
In [55]: # import SVC model library
```

```
from sklearn.svm import SVC
```

```
In [56]: svcModel = SVC(gamma='auto')
svcModel.fit(x_train, y_train)
```

```
Out[56]: SVC(gamma='auto')
```

```
In [57]: print('Accuracy on training data:', svcModel.score(x_train, y_train))
```

```
Accuracy on training data: 0.9775
```


```
In [58]: print('Accuracy on test data:', svcModel.score(x_test, y_test))
```

```
Accuracy on test data: 0.37
```

```
In [59]: # Visualize

plt.figure( figsize=(30,10))
#plt.xscale('log')
plt.hist([first500.Occupation,first500.Age], bins=10, label=['Occupation','Age',])
#plt.hist(combinedData.Age)
```

```
Out[59]: (array([[131., 124., 139., 106.,  0.,  0.,  0.,  0.,  0.,  0.],
                [ 53.,  0.,  0.,  0., 245.,  0., 131.,  0.,  0.,  71.]]),
array([ 1. ,  5.9, 10.8, 15.7, 20.6, 25.5, 30.4, 35.3, 40.2, 45.1, 50. ]),
<a list of 2 BarContainer objects>)
```



In []: